

The Ultimate Guide to Why Communities Fail

Twenty-four documented patterns and one the standard misses, what each looks like from inside, and the structure that prevents it – for people who can still change the rules.

6 LESSONS 36 MIN READ UPDATED 12 AUGUST 2026

Free to read, print, translate and argue with. Corrections are the contribution we would rather have than a payment.

ecohubs.community/learn

Contents

01	Power nobody voted for	6 min
	Five patterns that hand authority to people the group never chose – and why every one of them is built by somebody being helpful.	
02	Conflict that never resolves	5 min
	Two patterns, opposite in temperament and identical in effect – a community that cannot metabolise disagreement accumulates it.	
03	Where the money quietly moves	7 min
	Two communities met the same pressure and answered it opposite ways – one sold the commons and kept the model, the other kept the commons and changed the model.	
04	Who belongs, and who decides that	6 min
	Four patterns about the edges – what the community stands for, what it may ask of you, and how somebody stops being a member.	
05	The work that breaks people	6 min
	Five patterns that all end in the same place – a community running on two or three people who cannot stop, and a group that can see it and cannot fix it.	
06	Changes that cannot be undone	6 min
	Five patterns about how rules change – the temporary measure that stayed, the founding principle nobody protected, and the community that can no longer remember why.	

ABOUT THIS DOCUMENT

This is the whole guide as it stood on 12 August 2026. The quizzes are printed with their answers and explanations, which is what makes them useful on paper. The cost estimator on the web version is interactive; here you get the four equity models it is built on, and the accompanying spreadsheet does the arithmetic.

The live version, with working quizzes and the calculator, is at ecohubs.community/learn/guides/why-communities-fail.

Power nobody voted for

Five patterns that hand authority to people the group never chose – and why every one of them is built by somebody being helpful.

THE SHORT VERSION

Authority does not wait to be granted. If a community has not written down who decides what, the question gets answered anyway – by whoever has the most history, the most time, or the most stamina at eleven at night.

Five patterns do most of this work: **a clique that decides before the meeting**, **speakers who fill the room**, **a founder whose objection nobody tests**, **decisions with no stated type**, and **an authority that comes from how compelling somebody is**.

None of them is built by a villain. Every one is built by somebody being useful, and that is exactly why they are hard to name and harder to undo.

Start with the thing that makes this lesson difficult.

In almost every community where power has settled somewhere nobody chose, **the people holding it are working harder than anyone else**. They answer the emails at midnight. They remember why the rule exists. They are the reason the boiler works. Asking a group to constrain them feels like ingratitude, and it usually is experienced as ingratitude, which is why the conversation gets postponed for years.

So the framing that works is not *who has too much power*. It is **what happens when they leave** – and that question can be asked with warmth, out loud, in a meeting, without anybody being accused of anything.

The five, and what they have in common

They are separate patterns with separate fixes, but they share a mechanism: **a decision that was never assigned found a home**. Nobody took the authority. It was left on the table, and somebody sensible picked it up.

- **The clique that decides before the meeting** – the real conversation happens in the kitchen, and the meeting ratifies it
- **Speakers who fill the room** – nothing is blocked, but only some people's proposals ever get finished
- **The founder's informal veto** – no rule grants it, and everyone routes around it anyway
- **Decisions with no stated type** – the same question is operational or constitutional depending on who is asked
- **Authority from charisma** – a person is so compelling that disagreeing with them feels like a failure of understanding

If you recognise one, read the others before concluding you have found the whole problem. They co-occur, and they reinforce each other: a founder's unwritten veto is far more powerful in a group whose decisions have no stated type, because there is nothing to say the veto does not apply here.

Why it is hard to see from inside

Three reasons, and the third is the one that keeps communities stuck.

It looks like competence. Everything above is indistinguishable, day to day, from a group that has sensibly deferred to the person who knows most. The difference only shows up under stress – when the group wants to go a way that person does not.

Nobody experiences themselves as powerful. Ask the founder and they will tell you, sincerely, that they wish other people would take more on. That is usually true. Informal power is not a thing people feel themselves holding; it is a thing other people feel themselves working around.

And the group already knows. This is the part that surprises people. In our experience the pattern is almost never a secret – members will describe it accurately, in private, sometimes for years. What they lack is not perception. It is **a legitimate way to raise it that does not read as an attack on someone who has given a great deal.**

That is the whole argument for writing things down. Not because a written rule is wiser than a good person, but because a written rule can be discussed without anybody being the subject.

WHAT WE DON'T KNOW

Everything here is about structure, and structure is not the whole story. A community where one person is genuinely behaving badly has a conflict problem, not a governance problem, and no decision matrix will fix it – see **conflict that never resolves.**

The reverse mistake is more common though, and more expensive: treating a structural problem as a personal one. If the same argument keeps recurring with different people in the difficult role, it was never about the people.

If it is already happening

Stabilisation before prevention. A community in the middle of this does not need a governance redesign; it needs the next three months to be less bad.

Name the smallest true thing. Not “we have a power problem” – “I noticed the last four decisions about money were made between meetings.” A specific, checkable observation is arguable. A diagnosis is not, and it will be experienced as an accusation.

Ask the succession question, not the power question. *What happens to this if you are ill for a month?* It is a real question, it is kind, and it produces the same information.

Write down one decision. Not the whole matrix – one. Pick the category that keeps going wrong and state who decides it, by what method, within what limit. A community that has never written any of this down will find the first one takes an evening and the second takes twenty minutes.

What prevents it

A decision matrix

, and the thing communities forget to pair with it: a register of who currently holds each role. The matrix says what a role may do; the register says who is in it today, which is the question a new member actually has.

Then the two rules that make it hold:

- **A decision nobody can classify goes to the higher tier.** Without that, every borderline case drifts toward the easier meeting.
- **The table says how the table is amended.** A document that cannot describe its own amendment will be amended informally, by whoever notices first.

RCOS · LAYER 2

The RCOS Standard states the underlying rule without hedging: **authority must not be derived from charisma, seniority, ownership, or informal influence** (§4.3.4). Not *should not* – and it is the only clause in the layer that rules out something no community ever voted for.

Around it: all authority must be assigned to explicitly defined roles, circles or bodies (§4.3.1), carrying its scope, its limits and its term where it has one (§4.3.2); nobody may exercise authority outside their assigned scope (§4.3.3); and **decisions made outside the Decision Matrix are invalid** (§4.4.4).

That last one sounds severe until you notice what it protects. Without it, the fastest route to authority in any community is to keep acting outside the structure until acting outside the structure *is* the structure.

RCOS Core v0.1 – §4.3 Authority Boundaries

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Governance & Power** – the four documented patterns this lesson covers, in specification form
- 2 **Freeman, J. – The Tyranny of Structurelessness (1970/1972)**[↗] – why a structureless group is impossible, and what informal elites form around
- 3 **RCOS Core – Layer 2: Governance & Decision Logic** – authority boundaries, the decision matrix, and the safeguards against concentration
- 4 **Conflict, power and repair** – the joiner-facing version, for reading a community from the outside

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked.

How this is written

Conflict that never resolves

Two patterns, opposite in temperament and identical in effect – a community that cannot metabolise disagreement accumulates it.

THE SHORT VERSION

Only two documented patterns here, and they look like opposites. One community **never confronts anything** and calls it harmony. Another **reaches for consequences first** and calls it accountability.

They produce the same thing: disagreement that does not get processed, stored somewhere, and paid for later with interest.

The failure is never the conflict. It is the absence of a **route** – somewhere a disagreement can go that is not the meeting, the kitchen, or everyone's evening.

Every community will tell you it takes conflict seriously. The only useful question is **when the process was last used, by whom, and what changed afterwards** – and a group that has never triggered its process in ten years does not have a peaceful history. It has an unusable process.

The two

- **Conflict that is never named** – the group is warm, nothing is confronted, and difficulty is absorbed until somebody leaves
- **Punishment before repair** – the group does act, and reaches straight for the consequence, so nobody raises anything small

They are opposite temperaments and they arrive at the same destination, which is why a community fleeing one often lands in the other. A group exhausted by avoidance decides to “get serious about accountability” and discovers eighteen months later that nobody raises anything any more – because now raising something starts a process with a verdict at the end of it.

The stable version is neither. It is a **ladder**: several steps, the early ones cheap and reversible, escalation available and rarely needed.

Why the absence is invisible

A community with no conflict route does not experience itself as having a problem. It experiences a series of unrelated difficult situations, each with its own explanation – a personality clash, a stressful year, someone who was never quite right for the place.

The pattern only appears when you line them up and notice **the same shape**: raised privately, discussed informally, never concluded, and resolved eventually by somebody's departure. Departure is the default conflict process of any community that has not chosen one. It works, in the sense that the conflict ends. It costs a member each time.

WHAT WE DON'T KNOW

This lesson is about structure, and there is a limit to what structure reaches. **No process resolves a conflict where one party has substantially more power than the other and the community will not say so** – mediation between unequals produces agreement rather than resolution, and it resurfaces in eighteen months attached to something else.

That is the most misdiagnosed failure in this whole guide. The group concludes its process did not work and goes looking for a better method, when the problem was never the method. If that sounds familiar, **power nobody voted for** is the lesson you actually need.

If it is already happening

Sort what you have before you build anything. Most communities in trouble have three or four live difficulties and treat them as one atmosphere. They are not the same problem and they do not take the same route – a dispute about the rota is not a dispute about authority, and neither is a breach of a stated boundary.

The sort that matters most is: **is this actually between these two people?** Some conflicts are produced by a rule, an incentive or an allocation, and the two people arguing are simply standing where it pinches. Those survive mediation, because mediation fixes a relationship that was not broken.

Then take one, and take the smallest one. A community learning to run a repair process should not learn on its hardest case. Run the cheap version on something minor, find out where it breaks, and fix that.

What prevents it

Four questions, answerable in a single conversation, and the first one carries most of the weight:

- **Who may start it?** If only the board, the founders or a majority can trigger the process, it is unavailable to precisely the person most likely to need it – someone new, alone, and in conflict with somebody established.
- **What happens first?** Almost every workable system starts with the two people talking, with support.
- **Who holds it, and are they neutral?** In a community of thirty the obvious mediator is somebody's close friend. Small communities routinely need someone from outside, and the ones that have thought about it already know who to call.
- **What can it conclude?** A process that can only recommend is a conversation. Repair needs an available outcome: an agreement, a rule change, a consequence, an apology on the record.

RCOS · LAYER 4

The RCOS Standard's framing here is stronger than most communities would put it, and it is worth quoting because it removes the "we just haven't had much conflict" defence:

“Conflict MUST be treated as a handled condition with defined pathways; ignoring, suppressing, or normalizing unresolved conflict MUST be considered a system violation” (§6.6.1).

A violation – not a shortcoming, not a cultural preference. The absence of a route is itself the failure, regardless of whether anything has gone wrong yet.

Two more invariants sit beside it. **Repair and restoration must precede punishment except where immediate safety is at risk** (§6.6.3), which is the second pattern in this lesson stated as a rule. And **conflicts involving power asymmetries must trigger elevated safeguards** (§6.6.2) – the caveat above, made structural.

The resolution process must also be **accessible without requiring social status, seniority, charisma, or informal proximity to decision-makers** (§6.2.4). That is the “who may start it?” question, and RCOS answers it as an access requirement rather than a design preference.

RCOS Core v0.1 – §6.6 Layer Invariants

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Conflict & Accountability** – the two documented patterns this lesson covers, in specification form
- 2 **RCOS Core – Layer 4: Conflict, Repair & Accountability** – conflict classes, the resolution ladder, safeguards and proportional sanctions
- 3 **Ostrom, E. – Governing the Commons (1990)**⁷ – graduated sanctions and cheap, local conflict resolution, among the eight design principles
- 4 **Conflict & repair** – the five conflicts that actually recur, and what a repair process has to contain

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked. **How this is written**

Where the money quietly moves

Two communities met the same pressure and answered it opposite ways – one sold the commons and kept the model, the other kept the commons and changed the model.

THE SHORT VERSION

Communities rarely run out of money in a year. They run out over thirty, and every individual decision along the way was reasonable.

Three patterns: **the commons sold to cover operating costs**, **finances nobody can actually inspect**, and **contribution that quietly stops being shared**. All three are slow, all three are disclosed if anyone looks, and all three are invisible while they are working.

The most useful evidence in this whole guide sits here – two communities that met the same pressure and answered it in opposite directions, both of which published what they did.

Most writing about community money is about **how much**. Entry, monthly, exit. That matters, and it is **covered elsewhere**.

This lesson is about something slower: what happens to a community's economy over decades, when income does not quite cover costs and nobody has to decide anything dramatic. There is always a way to close this year's gap. The question is what closing it repeatedly does to the twenty-fifth year.

The three

- **The commons sold to cover the shortfall** – each sale funds another year, and the thing being sold is the thing that made the community possible
- **Finances nobody can inspect** – not hidden, exactly; just held by two people and never actually looked at
- **Contribution that stops being shared** – some members pay their way out of the work, and the arrangement is fine until it is the norm

Two communities, one pressure, opposite answers

This is the part of the guide with the strongest evidence behind it, and both cases are documented by the communities themselves.

The Findhorn Foundation is a registered Scottish charity, so it files audited accounts. In July 2023 it announced its configuration was no longer financially viable; educational programmes ceased after 22 September, and the trustees' report records that *"During October and November, 48 staff left the Trust*

on a redundancy basis, with a core team of 10 staff remaining over winter 2023 and spring 2024.” The report names Brexit, Covid, fires that destroyed two community facilities, and energy costs.

But the sentence that matters is not about any of those:

““We have for over 30 years sold non-core assets to enable the Foundation’s work to continue in spite of financial deficits, but the scale of our recent and forecast losses means that this is no longer a viable option.””

Thirty years, disclosed by the organisation itself, in a document an auditor signed. Alongside it: “The Foundation’s financial position has been weak for many years.”

The Farm, in Tennessee, hit the same wall in the early 1980s: by its own account it was spending “\$10,000 a week” against income of “around \$6,000”, having taken on “\$100,000 in debt overnight” when a crop froze in Florida, plus hospital bills over \$100,000.

Its answer went the other way. In 1983 it abandoned full communalism. Businesses were “*privatized, owned by their principal managers*” and had to start paying their employees; each adult took on “\$100 per person, plus an additional \$35 a month per person to go towards paying down our debt.” **The land and the common assets stayed common.** Within four years, by its own account, the community was debt free.

THE COMPARISON, STATED CAREFULLY

Same pressure – an income model that stopped covering costs. Opposite variables.

Findhorn kept the model and sold the assets. The Farm kept the assets and changed the model.

Only one of them still has the thing it was protecting. That is not a verdict on anybody: Findhorn published its own position honestly and restructured deliberately rather than collapsing, which the hub elsewhere calls **an ending that is not a failure**. And the Farm’s population fell from around 1,500 in 1979 to roughly 200 – with many of those departures, per the Tennessee Encyclopedia, coming *before* the restructuring rather than because of it.

What the pair shows is not which community was right. It is that **the variable is which one you treat as fixed**, and that most communities never notice they are choosing.

WHAT WE DON’T KNOW

One complication worth keeping, because the tidy version is wrong. Findhorn did not only sell. The same accounts record the Sanctuary project being **transferred to Park Ecovillage Trust**, trustees resolving to **gift Wikies Wood to Findhorn Hinterland Trust**, and continuing discussions about “*selling non-core assets into the community*.”

So part of the commons was **rehomed into other community bodies** rather than privatised out. Reading this as a straight sell-off would be the easy version and the inaccurate one.

Why it is hard to see

Because the annual decision is always defensible. No year's asset sale is the mistake. The mistake is the pattern, and a pattern is not a thing that appears on any agenda.

Because the alternative is worse in the short run. The community that refuses to sell has to cut something people care about, this year, in front of everyone. The community that sells has a slightly smaller commons and an unchanged programme. One of those conversations is much harder to have.

And because the accounting is honest. Findhorn's position was in its audited accounts the whole time. The information was never hidden. It was simply not the kind of information anyone reads as a trajectory.

If it is already happening

Work out the run rate, once. Not the budget – the gap. How much of last year's costs were covered by something that will not recur: a sale, a legacy, a grant, a reserve draw. That single number is the diagnosis, and most communities have never calculated it.

Then ask how many of those you have left. A community funding a €40,000 annual shortfall from asset sales, with three saleable assets, has three years. Said aloud, that changes a conversation which has been circling for a decade.

Separate the two decisions. *Can we afford this model?* and *what should we sell?* get answered together and should not be. Answering the second first is how thirty years happens.

What prevents it

A reserve fund with a stated purpose, and a rule that operating shortfalls may not be funded from capital without a decision at the constitutional tier. The point is not to forbid selling. It is to make selling a **decision** rather than a default – which is precisely what it stops being when it happens every year.

RCOS · LAYER 3

RCOS's Layer 3 exists to make economic power *"explicit, bounded, reviewable, and subordinate to governance"* – that last word being the one that does the work here.

Every resource in the governed scope must be classified as **commons or private** in a versioned registry, recording its steward, its access rules and **any constraint on transfer, sale or privatisation** (§5.1.1–2). That final field is what a thirty-year pattern needs in order to be visible: a sale constraint written down is a thing a decision has to override.

Two invariants close it. **"Resources declared as commons MUST NOT be privatized through informal, implicit, or unilateral action"** (§5.6.2) – and repetition is a form of implicit. And shared resources, flows and obligations **must be visible to the community by default**, with only limited and explicit exceptions (§5.6.1).

The treasury clauses add the missing half: spending authority must be explicitly bounded by thresholds and approval paths (§5.3.3), and the community must define **reserve, risk and liability policies including limits on debt and long-term obligations** (§5.3.6).

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Economy & Resources** – the three documented patterns this lesson covers, in specification form
 - 2 **[Findhorn Foundation – Trustees' Annual Report, year ended 31 March 2024](#)**[↗] – audited accounts, Scottish Charity SC007233 – the source of every Findhorn quotation here
 - 3 **[The Farm – Why The Farm Survived](#)**[↗] – the community's own account of the 1983 Changeover, and the figures quoted here
 - 4 **[The Farm – Tennessee Encyclopedia](#)**[↗] – the cooperative structure, the population figures, and that departures preceded the change
 - 5 **RCOS Core – Layer 3: Economic & Resource System** – the commons register, treasury transparency, and limits on accumulation
-

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked.

How this is written

Who belongs, and who decides that

Four patterns about the edges – what the community stands for, what it may ask of you, and how somebody stops being a member.

THE SHORT VERSION

Every community has a boundary. Most have never written down where it is, which means it gets located during an argument, about a specific person, by whoever is most upset.

Four patterns: **the purpose that quietly changed**, **the norm nobody could enforce**, **the rules that spread into private life**, and **removal with no process**.

They share a single missing thing – a stated answer to *what does this place actually require, and of whom*. Without it, the community defaults to the two worst options: enforcing nothing, or enforcing everything at once when it finally runs out of patience.

The edge of a community is where its structure is tested, and almost nothing about it is written down in advance.

That is not carelessness. Writing down what would get somebody removed feels, in year one, like planning for a betrayal that has not happened among people you have just chosen. So the question is deferred, and it stays deferred right up until the group is holding a live situation involving a real person with friends in the room.

The four

- **The purpose that quietly changed** – nobody voted to become something else, and enough people arrived wanting something else that it happened
- **The norm nobody could enforce** – a commitment central to the community's identity, with no process behind it
- **Rules that spread into private life** – the opposite failure, and the one that makes members leave without complaining
- **Removal with no process** – grounds invented during the case they are used on

Notice that two of these point in opposite directions. A community can fail by enforcing nothing and fail by enforcing too much, and **the same missing document causes both**. A group with no stated scope has no basis for acting and no basis for restraint, so it oscillates: years of tolerating everything, then one decisive over-correction that costs it three members.

The one that is missing from this list

There is a fifth pattern here that RCOS's stress tests do not catalogue, and it is worth naming because it decides how much every other rule on this page matters: **exit that is not really available**.

Nobody writes a rule forbidding departure. What blocks it is structural – an individual stake in the property that only converts to money when a buyer appears, an open-ended notice period, a queue with no position in it. Each part was put there by somebody being reasonable, and together they mean that leaving costs more than staying and disagreeing.

This is documented outside our field. Ran Abramitzky's study of the Israeli kibbutzim found that **a community's wealth functions as a "lock-in device" that increases the value of staying**, and that higher wealth reduces exit while sustaining a high degree of income equality. A stake you forfeit by leaving is a lock-in device whether or not anyone designed it as one.

It matters here because **every membership rule on this page is only as strong as the exit behind it**. A community whose members cannot leave has not achieved commitment. It has achieved captivity with good intentions, and its agreements are no longer consented to in any meaningful sense – they are merely endured.

WHAT WE DON'T KNOW

Being unable to leave and being unwilling to leave look identical from outside, and they feel similar from inside. The honest test is not whether anyone has left recently. It is **what leaving would cost, calculated in advance**, and whether the member could tell you that number today.

If it is already happening

Write the scope before you write the rules. A community arguing about whether it may have an opinion on somebody's diet, or their guests, or their income, is arguing about scope with only a specific case in front of it – which is the worst possible moment. What resolves it is a statement of what the community governs and, crucially, what it does not.

Handle the live case with what you have, and say so. If somebody has done something the group cannot tolerate and there is no process, the honest move is a process invented transparently for this case, stated as such, and reviewed afterwards. That is far better than pretending a rule existed, and far better than doing nothing.

Then write the general version afterwards, not during. Rules written in the middle of a case are written about a person, and everybody knows it – including the next person they get applied to.

What prevents it

Three documents most communities have none of: a **scope statement** saying what is governed and what is not, **identity constraints** stated so they could actually be tested, and an **exit and separation protocol** covering both directions – voluntary and forced.

RCOS puts the boundary in Layer 0 – constitutional, not operational, which is itself the argument. A community must **explicitly declare the scope of what it governs**, listing assets, decision domains and activities under collective control (§2.2.1–2), and must **explicitly list what is out of scope** (§2.2.3).

Then the rule that resolves the oscillation described above: **anything not explicitly declared as in scope must be treated as out of scope** (§2.2.4), and the community **must not exercise authority over persons, assets or domains that are out of scope** (§2.2.5). Silence resolves in the member's favour – which is the opposite of how an unwritten arrangement drifts.

On norms: a community must declare identity-level constraints that materially affect participation or behaviour (§2.4.1), and those constraints must be **testable and enforceable through defined processes** (§2.4.3) and **must not be enforced implicitly or informally** (§2.4.4). A value with no process behind it is not an identity constraint under RCOS. It is an aspiration, and the standard says so.

And underneath all of it, from Layer 1: **voluntary exit must be possible at all times** (§3.6.1), with exit procedures that are explicit, documented and **non-punitive** (§3.6.2), and **asset, role and responsibility separation defined before exit** rather than negotiated during it (§3.6.5).

RCOS Core v0.1 – §2.2 Scope Declaration, §2.4 Identity Constraints

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Membership & Boundaries** – the four documented patterns this lesson covers, in specification form
- 2 **RCOS Core – Layer 0: Identity & Scope** – purpose, scope, and identity constraints that have to be testable
- 3 **RCOS Core – Layer 1: Membership System** – membership states, rights and obligations, and the rules on exit
- 4 **Abramitzky, R. – The Limits of Equality: Insights from the Israeli Kibbutz (QJE, 2008)⁷** – community wealth as a lock-in device that raises the value of staying
- 5 **Membership: getting in, and getting out** – the joiner-facing version – the four stages and what a fair agreement contains

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked. **How this is written**

The work that breaks people

Five patterns that all end in the same place – a community running on two or three people who cannot stop, and a group that can see it and cannot fix it.

THE SHORT VERSION

The largest pattern in this guide, and the one communities are best at describing and worst at fixing.

Five modes: **care work nobody counts**, **responsibility that quietly becomes authority**, **growth outpacing onboarding**, **self-sufficiency without contribution**, and **information nobody can reach**.

They converge. A community running on two or three indispensable people has all five at once, and the reason it stays that way is not ignorance – it is that **redistribution needs a structure, and complaining is not one**.

Here is the thing that makes this lesson different from the others.

With power, with conflict, with money, the community usually cannot see the pattern. **With this one it can.** Members will describe it accurately, in detail, unprompted. They will name the person carrying too much. They will say it is not sustainable. And then nothing will change, for years.

That gap – between accurate diagnosis and no repair – is the whole subject.

The five

- **Care work nobody counts** – the noticing, remembering and smoothing that never reaches a rota
- **Responsibility that becomes authority** – the person who does it ends up deciding it, without anybody granting that
- **Growth that outpaces onboarding** – new members arrive faster than the community can teach itself to them
- **Self-sufficiency without contribution** – households quietly withdraw into their own lives, and the commons is maintained by whoever minds
- **Information nobody can reach** – the accounts, the passwords, the history, the reason the rule exists

Why knowing is not enough

FROM A COMMUNITY WE WORKED WITH

One person knew everything and held nearly every responsibility. They were visibly overwhelmed. Other members said so – openly, repeatedly, in the community's own words: he has too much.

And nobody found a way to spread it.

Not because they did not care, and not because he refused. Because “you have too much” is a description, and a description cannot be acted on. What was needed was a strategic approach – somebody sitting down with the list of what he actually did, splitting it into named pieces, and finding a holder and a handover for each. Nobody did that, because it is a project, and projects need an owner, and the only person who understood the list well enough to own it was him.

That is the mechanism, and it generalises. **The people best placed to redistribute the work are the people currently doing it, and they are the people with no capacity to run a redistribution project.** So the group circles: it notices, it says so, it feels bad, and the situation persists – while everyone involved behaves reasonably and cares about each other.

Three things follow, and they are the practical content of this lesson:

Complaining is not an intervention. A community that has raised something twenty times and seen no change does not have a communication problem. It has a missing structure, and more conversation will not produce one.

The work has to be enumerated before it can be shared. You cannot hand over “the admin”. You can hand over “renewing the insurance in March”, “reconciling the account quarterly”, and “keeping the key register” – three named things with three holders. The enumeration is most of the labour, and it is the step that gets skipped.

Somebody who is not overloaded has to run it. This is the part communities get wrong most reliably. Asking the overloaded person to organise their own unloading is asking for one more task from the person with the least room, and it is why the attempt fails on the third week every time.

If it is already happening

Do the audit, not the conversation. One person, two hours, a list: everything that has to happen for this community to keep functioning, and who currently does it. Not who is *supposed* to. Who does.

Look for the single points of failure. Anything with one name against it and no second person who could do it. That list is usually shorter than people fear and more alarming than they expect.

Then take the top three and give each one a name, a scope and a handover date. Not all of it. Three. A community that redistributes three responsibilities properly has learned how, and can do the rest; a community that tries to redistribute everything produces a spreadsheet nobody maintains.

And write the limit down. How many roles one person may hold, or how many hours a month is too many. Without a stated limit there is no moment at which anybody is entitled to say *this is too much* without it sounding like a complaint about themselves – which is exactly why the overloaded person never says it.

What prevents it

Named roles with terms and handovers, a stated ceiling on how much any one person carries, and a definition of contribution that counts the **invisible work** as work. None of that is about fairness. It is about a community being able to continue when somebody is ill.

RCOS states the enumeration step as a requirement, which is unusual and is the point: “**All ongoing responsibilities MUST be assigned to explicit, named roles rather than implicit expectations or informal agreements**” (§7.1.1). And its protective mirror – **no person may be held accountable for responsibilities not formally assigned to a role** (§7.1.4).

Every role enters a Role Registry with its purpose, decision authority, boundaries, and its appointment, review and removal process; term length and rotation are recorded where a role has them (§7.1.2). Each role must also carry an accountability mechanism defining **how underperformance, overload, or role failure is handled, and how handover and knowledge transfer occur** (§7.1.3) – overload named as a role-design question rather than a personal one.

Workload boundaries are a required artifact, not a courtesy: explicit limits on meeting load and role load, and a named mechanism for **relief, substitution or redistribution** (§7.4.2). And the trigger clause: **persistent overload, burnout risk, chronic non-participation, or dependency on over-functioning individuals MUST trigger review or repair processes** as defined in Layer 4 (§7.4.4).

Then the clause that names the case above directly: “**The community MUST ensure that no single individual is a critical single point of failure for core operations**” (§7.5.1), with documented procedures, clear handover mechanisms and backup arrangements where feasible (§7.5.2).

Layer 3 closes it from the economic side: a community **must not structurally depend on unpaid, invisible or informal labour for system survival** without defining the corresponding recognition, obligation or compensation (§5.2.3).

RCOS Core v0.1 – §7.1 Roles and Responsibilities, §7.4 Workload and Capacity Boundaries

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Operations & Coordination** – the five documented patterns this lesson covers, in specification form
- 2 **RCOS Core – Layer 5: Operations & Coordination** – the Role Registry, workload boundaries, and the single-point-of-failure clause
- 3 **Freeman, J. – The Tyranny of Structurelessness (1970/1972)**⁷ – delegate specific authority for specific tasks, rotate them, and spread information widely
- 4 **Daily life in community** – the work nobody put on the rota, and what a week actually contains

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked. **How this is written**

Changes that cannot be undone

Five patterns about how rules change – the temporary measure that stayed, the founding principle nobody protected, and the community that can no longer remember why.

THE SHORT VERSION

A community that cannot change its rules dies of rigidity. One that changes them without noticing dies of drift. Both are failures of the same missing thing: **a stated way for a rule to change, and a record of when it did.**

Five patterns: **the invariant nobody protected, the emergency measure that stayed, governance with no path to change it, experiments with no way back, and a community that forgot why.**

The defence is not vigilance. It is bookkeeping – dating the agreements and keeping the old versions – which is unglamorous, cheap, and the only thing that works.

Every rule in a community was written by people responding to a situation. Ten years later some of those situations are gone, some of the rules are load-bearing, and almost nobody can tell which is which.

That is the subject of this lesson: **not what the rules are, but how they move** – and the fact that most communities have no account of their own movement at all.

The five

- **The invariant nobody protected** – the founding principle everybody assumed was permanent, changeable by an ordinary majority on a Tuesday
- **The emergency measure that stayed** – decided fast because it had to be, never reviewed, now precedent
- **Governance with no path to change it** – the opposite failure, and the one that produces departures rather than arguments
- **Experiments with no way back** – a trial that quietly became the arrangement, because nobody set an end date
- **A community that forgot why** – the rule survives, the reason does not, and nobody can evaluate either

And one more that RCOS does not catalogue: **sub-groups with undefined autonomy**, which is what happens when a community grows into parts without deciding what the parts may decide.

Why drift is the hardest of them

Four of these announce themselves eventually. Drift does not, and it is the one most communities are actually living in.

Nobody notices a slow forgetting by paying closer attention, because **there is no moment to notice**. The agreements stop being read, then stop being taught, then describe a place that no longer exists – and every step of that is a non-event.

Which is why vigilance is the wrong defence and bookkeeping is the right one. A community that can say **which version is in force, when it last changed, and who decided it** has turned a feeling into a check anybody can run.

The difference that makes is mostly social. Without a current text, raising drift means saying *I think we have lost something* – a criticism of everyone present, and heard as one. With it, the same observation is that the practice and the document disagree. A discrepancy, pointed at rather than levelled, and the sort of thing a group can look into without anyone having to be at fault.

FROM A COMMUNITY WE WORKED WITH

A founder produced a new version of a foundational document. Nobody could say whether it had ever passed the community's consensus process – the original was old enough that the memory had gone.

So it sat alongside the previous version, with a new version number, **no record of adoption, and no statement of what had changed between them**.

The community was not in conflict about the content. It could not establish **which of its own documents was in force**, or what the difference was. Everything downstream of that – every rule, every decision citing it – inherited the ambiguity.

Note how little would have prevented it: an adoption record and a summary of changes. Two fields.

If it is already happening

Establish the present before you reconstruct the past. Which documents are in force, as of today, agreed by whoever is in the room. Not their history – their status. This can be done in one meeting and it stops the bleeding.

Then date everything from now on. A change record with four fields: what changed, who decided it, under which rule, and from when. A community with two clean years and a murky past is in a far better position than one that spent six months excavating and gave up.

Keep the superseded versions. Deleting the old text destroys the only evidence of what the community used to think, which is what you need when somebody asks why the rule exists.

And review the emergency measures specifically. Every community has two or three arrangements adopted quickly during something difficult and never revisited. List them, and give each one a date by which it is either ratified properly or ends.

What prevents it

A **change protocol** that says how a rule changes and who decides, a **version history** that makes “which version is in force” answerable, and a **learning log** so that what a difficult year taught does not leave with the people who lived it.

All three are boring. That is the point – this is the one failure class where the defence is entirely clerical, and a community that finds that beneath it will forget why it exists.

RCOS · LAYER 6

RCOS's sixth layer treats evolution as a governed process rather than improvisation, and its central clause is a validity rule: “**No informal, undocumented, or ‘understood’ rule changes MAY be considered valid**” (§8.2.5).

Every adopted change must be versioned and traceable (§8.2.1), and the **Version History** must record the version identifier, adoption and effective dates, the **decision record reference – authority, mechanism, threshold** – a summary of changes, and any migration notes (§8.2.2). At any point the community must be able to determine **which version is currently in force and which artifacts are authoritative** (§8.2.3), and superseded rules must remain accessible **together with the dates during which they were in force** (§8.2.4).

The layer invariants set the two boundaries this lesson runs between: **change must be possible but constrained – no change may be instantaneous, implicit, or unreviewable** (§8.7.1), and **major failures and adaptations must be captured as shared learning, not erased or hidden** (§8.7.4).

Underneath it, from Layer 0: **no decision, role, process, or emergency measure may override an invariant** (§2.3.4), and where an invariant conflicts with any other rule, **the invariant prevails** (§2.3.5).

RCOS Core v0.1 – §8.2 Versioning and Authority, §8.7 Layer Invariants

SOURCES & FURTHER READING

- 1 **RCOS stress tests – Change & Emergencies** – the five documented patterns this lesson covers, in specification form
- 2 **RCOS Core – Layer 6: Evolution & Adaptation** – change mechanisms, version history, time-bounded experiments and the learning log
- 3 **Just For Now: how emergency rules become permanent** – the longer EcoHubs argument about temporary measures that outlive their emergency
- 4 **Resilience & why communities fail** – drift as one of the six recurring patterns, and why it is the only one that does not feel like anything

Written by EcoHubs members, with AI assistance for drafting and editing, and reviewed by a person before publication. Facts are checked against the sources listed; anything we could not verify is marked.

How this is written